

Hadeel Abdeljalil

## HDFS commands:

```
hadoop fs -mkdir /user/cloudera/sportshop
```

```
hadoop fs -mkdir /user/cloudera/sportshop/input
```

I copied the data provided and added it to a file called salesData.txt to /home/cloudera =>file path

```
hadoop fs -put salesData.txt /user/cloudera/sportshop/input/
```

```
hadoop fs -cat /user/cloudera/sportshop/input/salesData.txt
```

## Using HDFS commands, show to do the following:

- run the job

```
hadoop jar test.jar test. SalesAnalysis /user/cloudera/sportshop/input /user/cloudera/sportshop/output
```

Note: I added SalesAnalysis at test package

- browse the content of 'output' directory

```
hadoop fs -ls /user/cloudera/sportshop
```

- show the content of the output files

```
Hadoop fs -cat /user/cloudera/sportshop/output
```

## the full code of the MapReduce job:

```
import org.apache.hadoop.conf.Configuration;

import org.apache.hadoop.fs.Path;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
```

```
import java.io.IOException;
```

```
public class SalesAnalysis {
```

```
    // First Job Mapper
```

```
    public static class SalesMapper extends Mapper<Object, Text, Text, LongWritable> {
```

```
        private final static LongWritable totalSalesAmount = new LongWritable();
```

```
        private Text retailerCity = new Text();
```

```
        @Override
```

```
        public void map(Object key, Text value, Context context) throws IOException, InterruptedException {
```

```
            if (value.toString().startsWith("Retailer")) {
```

```
                return;
```

```
            }
```

```
            String[] fields = value.toString().split(",");
```

```
            if (fields.length == 6) {
```

```
                try {
```

```
                    String retailer = fields[0].trim();
```

```
                    String city = fields[2].trim();
```

```
                    int pricePerUnit = Integer.parseInt(fields[4].replace("$", "").replace(",", "").trim());
```

```
                    int unitsSold = Integer.parseInt(fields[5].replace(",", "").trim());
```

```
                    long salesAmount = (long) pricePerUnit * unitsSold;
```

```
                    retailerCity.set(retailer + ", " + city);
```

```
                    totalSalesAmount.set(salesAmount);
```

```
                    context.write(retailerCity, totalSalesAmount);
```

```
                } catch (NumberFormatException e) {
```

```
                    System.err.println("Error parsing record: " + value.toString() + " " + e.getMessage());
```

```
    }  
    }  
}  
}
```

// First Job Reducer

```
public static class SalesReducer extends Reducer<Text, LongWritable, Text, LongWritable> {  
    private LongWritable result = new LongWritable();  
  
    @Override  
    public void reduce(Text key, Iterable<LongWritable> values, Context context) throws IOException,  
        InterruptedException {  
        long sum = 0;  
        for (LongWritable val : values) {  
            sum += val.get();  
        }  
        result.set(sum);  
        context.write(key, result);  
    }  
}
```

// Second Job Mapper

```
public static class SortMapper extends Mapper<Object, Text, LongWritable, Text> {  
    private Text retailerCity = new Text();  
    private LongWritable salesAmount = new LongWritable();  
  
    @Override  
    public void map(Object key, Text value, Context context) throws IOException, InterruptedException {  
        String[] fields = value.toString().split("\\t");  
        if (fields.length == 2) {  
            retailerCity.set(fields[0]);  
        }  
    }  
}
```

```

        salesAmount.set(Long.parseLong(fields[1]));

        // The negative sign makes the LongWritable value descending
        context.write(new LongWritable(-salesAmount.get()), retailerCity);
    }
}

// Second Job Reducer

public static class SortReducer extends Reducer<LongWritable, Text, Text, LongWritable> {

    @Override

    public void reduce(LongWritable key, Iterable<Text> values, Context context) throws IOException,
    InterruptedException {

        for (Text val : values) {

            // Negate the key back to its original positive value

            context.write(val, new LongWritable(-key.get()));

        }

    }

}

public static void main(String[] args) throws Exception {

    // Configuration and First Job

    Configuration conf = new Configuration();

    Job job1 = Job.getInstance(conf, "sales sum");

    job1.setJarByClass(SalesAnalysis.class);

    job1.setMapperClass(SalesMapper.class);

    job1.setCombinerClass(SalesReducer.class);

    job1.setReducerClass(SalesReducer.class);

    job1.setOutputKeyClass(Text.class);

```

```
job1.setOutputValueClass(LongWritable.class);

FileInputFormat.addInputPath(job1, new Path(args[0]));
FileOutputFormat.setOutputPath(job1, new Path(args[1] + "/temp"));

boolean success = job1.waitForCompletion(true);
if (!success) {
    System.exit(1);
}

// Configuration and Second Job
Job job2 = Job.getInstance(conf, "sort by sales amount");

job2.setJarByClass(SalesAnalysis.class);
job2.setMapperClass(SortMapper.class);
job2.setReducerClass(SortReducer.class);

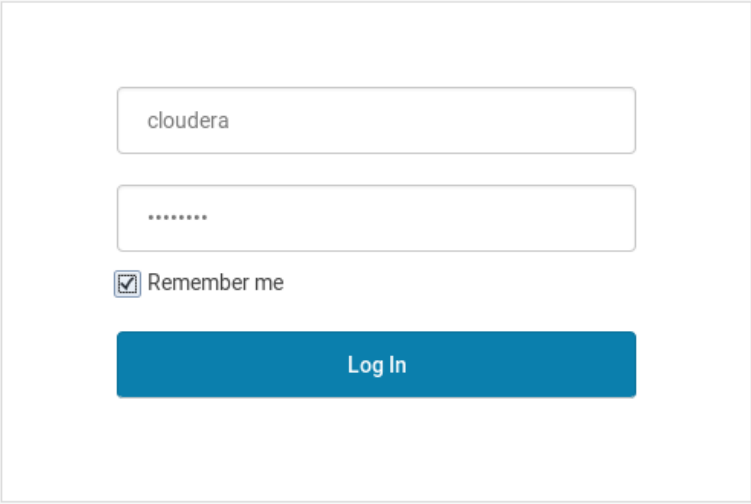
job2.setOutputKeyClass(LongWritable.class);
job2.setOutputValueClass(Text.class);

FileInputFormat.addInputPath(job2, new Path(args[1] + "/temp"));
FileOutputFormat.setOutputPath(job2, new Path(args[1] + "/final"));

System.exit(job2.waitForCompletion(true) ? 0 : 1);
}
}
```

## Steps:

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
[QuickStart] Shutting down CDH services via init scripts...  
kafka-server: unrecognized service  
[QuickStart] Disabling CDH services on boot...  
error reading information on service kafka-server: No such file or directory  
[QuickStart] Starting Cloudera Manager server...  
[QuickStart] Waiting for Cloudera Manager API...  
[QuickStart] Starting Cloudera Manager agent...  
[QuickStart] Configuring deployment...  
Submitted jobs: 81  
[QuickStart] Deploying client configuration...  
Submitted jobs: 82  
[QuickStart] Starting Cloudera Management Service...  
Submitted jobs: 90  
[QuickStart] Enabling Cloudera Manager daemons on boot...  
  
Success! You can now log into Cloudera Manager from the QuickStart VM's browser:  
  
http://quickstart.cloudera:7180  
  
Username: cloudera  
Password: cloudera  
[cloudera@quickstart ~]$
```



A login form for Cloudera Manager. It features a text input field containing the username 'cloudera', a password input field with masked characters '\*\*\*\*\*', a checked checkbox labeled 'Remember me', and a blue 'Log In' button.

Start HDFS & yarn(mapReduce), check health\*


## Overview 'quickstart.cloudera:8020' (active)

Started:	Sat Aug 24 07:37:42 -0700 2024
Version:	2.6.0-cdh5.13.0, r42e8860b182e55321bd5f5605264da4adc8882be
Compiled:	Wed Oct 04 11:08:00 -0700 2017 by jenkins from Unknown
Cluster ID:	CID-a24185f9-a545-40fe-9553-84c3fdca489f
Block Pool ID:	BP-1067413441-127.0.0.1-1508775264580

## Summary

Security is off.  
Safemode is off.

Block Pool Used:	832.38 MB (1.49%)
DataNodes usages% (Min/Median/Max/stdDev):	1.49% / 1.49% / 1.49% / 0.00%
Live Nodes	1 (Decommissioned: 0, In Maintenance: 0)
Dead Nodes	0 (Decommissioned: 0, In Maintenance: 0)
Decommissioning Nodes	0
Entering Maintenance Nodes	0
Total Datanode Volume Failures	0 (0 B)
Number of Under-Replicated Blocks	0
Number of Blocks Pending Deletion	0
Block Deletion Start Time	Sat Aug 24 07:37:42 -0700 2024



Cluster

- About
- Nodes
- Applications
  - NEW
  - NEW SAVING
  - SUBMITTED
  - ACCEPTED
  - RUNNING
  - FINISHED
  - FAILED
  - KILLED
- Scheduler

Tools

All Applications

Logged in as: drwho

Cluster Metrics

Apps Submitted	0	Apps Pending	0	Apps Running	0	Apps Completed	0	Containers Running	0	Memory Used	0 B	Memory Total	3 GB	Memory Reserved	0 B	VCores Used	0	VCores Total	3	VCores Reserved	0
----------------	---	--------------	---	--------------	---	----------------	---	--------------------	---	-------------	-----	--------------	------	-----------------	-----	-------------	---	--------------	---	-----------------	---

Cluster Nodes Metrics

Active Nodes	1	Decommissioning Nodes	0	Decommissioned Nodes	0	Lost Nodes	0	Unhealthy Nodes	0	Rebooted Nodes	0
--------------	---	-----------------------	---	----------------------	---	------------	---	-----------------	---	----------------	---

User Metrics for drwho

Apps Submitted	0	Apps Pending	0	Apps Running	0	Apps Completed	0	Containers Running	0	Containers Pending	0	Containers Reserved	0	Memory Used	0 B	Memory Pending	0 B	Memory Reserved	0 B	VCores Used	0	VCores Pending	0	VCores Reserved	0
----------------	---	--------------	---	--------------	---	----------------	---	--------------------	---	--------------------	---	---------------------	---	-------------	-----	----------------	-----	-----------------	-----	-------------	---	----------------	---	-----------------	---

Show 20 entries

ID	User	Name	Application Type	Queue	StartTime	FinishTime	State	FinalStatus	Running Containers	Allocated CPU VCoers	Allocated Memory MB	Reserved CPU VCoers	Reserved Memory MB	Progress	Tracking UI
No data available in table															

Showing 0 to 0 of 0 entries



# All Applications

## Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	VCores Used	VCores Total	VCores Reserved
2	0	0	2	0	0 B	3 GB	0 B	0	3	0

## Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes	Rebooted Nodes
1	0	0	0	0	0

## User Metrics for dr.who

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Containers Pending	Containers Reserved	Memory Used	Memory Pending	Memory Reserved	VCores Used	VCores Pending	VCores Reserved
0	0	0	0	0	0	0	0 B	0 B	0 B	0	0	0

Show 20 entries
 

Search:

ID	User	Name	Application Type	Queue	StartTime	FinishTime	State	FinalStatus	Running Containers	Allocated CPU VCoers	Allocated Memory MB	Reserved CPU VCoers	Reserved Memory MB	Progress	Tracking UI
<a href="#">application_1724510262132_0002</a>	cloudera	sort by sales amount	MAPREDUCE	root.users.cloudera	Sat Aug 24 07:52:19 -0700 2024	Sat Aug 24 07:53:22 -0700 2024	FINISHED	SUCCEEDED	N/A	N/A	N/A	N/A	N/A	<div></div>	<a href="#">History</a>
<a href="#">application_1724510262132_0001</a>	cloudera	sales sum	MAPREDUCE	root.users.cloudera	Sat Aug 24 07:49:55 -0700 2024	Sat Aug 24 07:52:17 -0700 2024	FINISHED	SUCCEEDED	N/A	N/A	N/A	N/A	N/A	<div></div>	<a href="#">History</a>

Showing 1 to 2 of 2 entries
 

First
 Previous
 1
 Next
 Last